

AIUC-1

# 2026: The End of Vibe Adoption

Executive briefing on AI security & latest research

Presented by Sanmi Koyejo (Stanford Trustworthy AI Research)  
& the AIUC-1 Consortium

## Co-authors

**Sanmi Koyejo**, Leader, Stanford Trustworthy AI Research Lab; Co-founder, VirtueAI

**Adnan Dakhwe**, CISO & CIO, DelphinusCyber

**Amin Jan**, Chief AI Architect, Department of War

**Brad Arkin**, Security, AIUC-1

**Brett Cumming**, Fortune 500 CISO

**Brian Levine**, Founder & Executive Director, FormerGov.com

**Cassie Crossley**, Fmr. VP, Supply Chain Security - Schneider Electric

**Chris DeNoia**, Founding Member, Tuskira AI Security Council (AISC)

**Chris Kirschke**, Field CISO, Tuskira

**Chris Monson**, Trustworthy AI Lead, Atlassian

**Chris Sandulow**, CISO, Confluent

**Christian Gorke**, CISO, Deutsche Börse

**Craig Weatherhead**, SVP IT Infrastructure & Security, Fastenal

**David Campbell**, Head of AI Security Research, Scale AI

**Dr. David Mussington**, Professor of Practice, University of Maryland School of Public Policy

**Gagandeep Singh**, VP, Global Compliance & Certification, Salesforce

**Jen Easterly**, CEO, RSAC

**Jonathan Fuller**, CISO, United States Military Academy

**Julie C. Chatman**, CISO, Human Health Project; CEO, ResilientTech Advisors

**Jyoti Wadhwa**, AI Governance & Enterprise Trust Executive

**Dr. Keri Pearlson**, Senior Lecturer & Principal Research Scientist, MIT Sloan

**Lena Smart**, Ambassador, AIUC-1

**Louise McElvogue**, Board Director and Advisor

**Mandy Andress**, CISO, Elastic

**Nancy Wang**, CTO, 1Password; Venture Partner, Felicis Ventures

**Neil Bennett**, CISO, Post Office Ltd.

**Omar Khawaja**, VP & Field CISO, Databricks

**Peter Holcomb**, Founder & CEO, Optimo IT

**Phil Venables**, Fmr. CISO, Google Cloud

**Rajiv Dattani**, Co-founder, AIUC

**Rohit Parchuri**, CISO & Advisor

**Rune Kvist**, Co-founder, AIUC

**Scott Kennedy**, CISO & DPO

**Scott Roberts**, CISO, UiPath

**Simon Goldsmith**, CISO

**Tim Mortimer**, Consulting Leader, Mandiant (Part of Google Cloud)

**Valentina Poghosyan**, Chief Compliance Officer, MongoDB

**Xabier Muruaga**, Global Head of AI & Data, Iberdrola

**Zachary Elewitz**, Fortune 500 AI Leader, Professor, and Board Member

**Edited by:** Emil Lassen, AIUC-1

*Thank you to Emil Pellonpaa, Abby Shen and Elle O'Hill for research support with this paper*

### About AIUC-1

AIUC-1 is the standard for AI agent security, safety, and reliability. The standard is developed with Technical Contributors from trusted institutions including Stanford, Orrick, MITRE, the Cloud Security Alliance, Google Cloud, Cisco, MongoDB, and more. Frontier AI agent builders pursue certification to demonstrate that their agents can be trusted. The full standard is available at [aiuc-1.com](https://aiuc-1.com)

## Executive Summary

**The era of vibe adoption is over.** After two years of experimental AI deployments driven by excitement and fear of missing out, enterprise security leaders are confronting a new reality: AI agents can now handle customer data, make business decisions, and integrate with critical systems - yet most organizations lack systematic ways to evaluate their safety and security.

**The risks are no longer theoretical.** [EY's recent survey](#) found 64% of companies with \$1bn+ turnover have lost over US\$1 million to AI failures. Incidents made headlines every week in 2025: chatbots hallucinating and offering refunds when it shouldn't, massive data leaks exposing millions of people's PII, successful jailbreak attempts.

### We believe 2026 will be defined by three challenges:

1. **The agent challenge:** AI has moved from "assistant" to "autonomous actor." Agents don't just suggest - they execute, making decisions and orchestrating tasks without step-by-step instructions. This introduces a new kind of risk: your own AI might take the wrong actions on its own, causing damage even without an attacker.


**The companies who will emerge as winners in 2026 recognize that high-level AI governance frameworks fall short in meeting this challenge.** Organizations need technically grounded frameworks that are AI agent specific to meet this new era and deploy AI agents with confidence

2. **The visibility challenge:** 63% of employees who used AI in 2025 pasted [sensitive company data into personal chatbots](#), including source code, customer data, and internal documents. Companies were stuck between banning AI and Wild West adoption, with 1 in 5 organizations reporting costly breaches due to "shadow AI". Security teams were left without the frameworks and tools needed to manage the challenge

**The companies who will emerge as winners in 2026 reframe AI security from compliance to competitive advantage.** Just like seatbelts enabled us to drive faster in a safe way, proactive AI governance will enable organizations to adopt AI with speed and confidence. Industry-led frameworks have emerged, empowering security & GRC leaders to guide their organizations through this challenge successfully.

3. **The trust challenge:** You can't trust your inputs. You can't trust your components. Prompt injection moved from research curiosity to production incidents in 2025 impacting organizations like Microsoft, Amazon, and ServiceNow. The foundations of secure system design - validate inputs, verify dependencies - break down when inputs are natural language and dependencies are opaque neural networks.

**The companies who will emerge as winners in 2026 require frequent technical testing of deployed agents.** Technical safeguards and good governance are sound foundations, but every new foundation model update, attack pattern, and system prompt tweak will result in different agent behavior. The only way to confidently deploy agents when customer trust and brand is on the line is by integrating agent red-teaming as a core tool.



**This executive briefing gives you the latest research findings from Stanford’s Trustworthy AI Research Lab**, led by Professor Sanmi Koyejo, combined with real world observations from executives in the AIUC-1 consortium who navigate these challenges every day. They span security, risk, and legal leaders across industry, government, academia, and non-profits.

**The briefing has 3 sections:** First, we outline why 2025 was an inflection point for AI adoption, drawing on the latest research coming out of Stanford. Second, we outline the 3 challenges we believe will define 2026, bridging research with real world incidents. Third, we summarise the win conditions for enterprises adopting AI securely in 2026.

“

The value proposition of agents is removing humans from the decision loop - but this also removes our most reliable safety mechanism.

**Professor Sanmi Koyejo**

Leader of Stanford’s Trustworthy AI Research Lab

## The 2025 inflection point

**In 2025, AI adoption outpaced AI security by every measure.** Incidents moved from CVE databases to weekly headlines in mainstream media. Air Canada faced a [lawsuit](#) after their AI agent hallucinated a refund policy. McDonald's hiring bot [leaked](#) millions of candidates' sensitive data. Salesforce's AgentForce was exposed to a critical [vulnerability](#), enabling attackers to exfiltrate sensitive CRM data through an indirect prompt injection attack.

The dominant shift in 2025 was not the emergence of new attack primitives but the collapse of friction required to operationalize them. Agentic tooling lowered the barrier to entry while simultaneously increasing systems coupling, creating failure modes that traditional controls were never designed to observe.

### What changed:

- OpenClaw, Cursor, Claude Code, and similar tools democratized agentic workflows, encouraging broader and faster adoption than security teams anticipated
- Incumbent SaaS vendors and new players added agents into and on top of existing platforms, expanding attack surfaces and enabling agent chaining
- Availability of agentic tools dramatically increased, barrier to entry decreased, quality of outcomes increased - all at once

This democratization of AI agents came at a cost. [EY's recent survey](#) found 64% of companies with \$1bn+ turnover have lost over US\$1 million to AI failures. One in five organizations [suffered a shadow AI breach](#), where unauthorized use of AI products exposed sensitive data.

The forces at play are strong. Nearly all businesses are requiring AI adoption at greater pace due to competitive pressure. Board members are asking pointed questions about AI risk. Regulators are publishing requirements. Insurance underwriters want documentation.

For CISOs, the question is no longer "should we adopt AI?" but "how do we ensure the AI we're adopting meets our standards?"

“

Agentic tools became faster to adopt, easier to use, and better at delivering results - all at once. The attack surface expanded just as quickly.

**Brett Cumming**

Fortune 500 CISO

## What research is showing

The latest research from Stanford's Trustworthy AI Research Lab points to these challenges only being magnified in 2026 as attacks become more sophisticated, with multistep and sequential attacks becoming the normal:

- **Multi-step attacks will become the norm.** Traditional security monitoring is built to detect discrete malicious events - a SQL injection, a privilege escalation, an unauthorized data access. Our research demonstrates that AI agents are vulnerable to STAC ([Sequential Task-level Attack Chains](#)) patterns that unfold over dozens or hundreds of steps, where no single step appears malicious in isolation. An attacker might subtly steer an agent through a seemingly benign conversation that gradually shifts its context, priorities, and eventual actions - making detection exponentially harder than traditional exploits.
- **AI systems are discovering novel attack patterns faster than humans can document them.** OpenAI's recent research revealed that their reinforcement learning-trained attacker could steer agents into harmful workflows spanning hundreds of steps, including "novel attack strategies that did not appear in human red teaming." This finding is critical: we can no longer rely solely on known attack patterns and human-curated defenses. The systems we're deploying can discover and execute attack vectors we haven't imagined, at speeds and scales that outpace traditional security response cycles.
- **The open web remains fundamentally unsafe for autonomous agents.** OpenAI has publicly acknowledged that AI browsers are vulnerable to prompt injection attacks - a tacit admission that agents cannot safely operate on untrusted web content without additional safeguards. This raises existential questions for agentic workflows: if an agent can be compromised simply by visiting a malicious website or reading untrusted content, how can organizations confidently deploy agents that interact with customers, browse documentation, or integrate with external systems?

## The end of vibe adoption

The next section of this briefing examines three interconnected problems that define AI security in 2026, drawing on additional research and insights from AIUC-1 consortium members navigating these risks on a daily basis.

A common thread connects all three challenges: traditional security controls - static policies, point-in-time assessments, human-centric access models - weren't designed for systems that change constantly, act autonomously, and process natural language as input.

Across these challenges, a consistent operational pattern emerges: AI introduces new execution dynamics that strain existing governance, identity, and data control models. Where ownership and control boundaries are unclear, autonomous behavior compounds small control weaknesses quickly.

The challenge for 2026 is clear: Vibe adoption is dead. Relying on existing security controls won't work. The organizations who fail to adopt new security measures in 2026 will continue to see costly AI failures damaging brand and customer trust.

“

As AI transitions from simple assistants to autonomous actors, traditional security controls are no longer sufficient. Organizations must move beyond high-level governance and adopt technically grounded frameworks to turn AI security into a competitive advantage.

**David Campbell**

Head of AI Security Research - Scale AI

# Three challenges defining AI security in 2026

## 1. The agent challenge

**AI has moved from "assistant" to "autonomous actor."** Agents don't just suggest - they execute, making decisions and orchestrating tasks without step-by-step instructions. This introduces a new kind of risk: your own AI might take the wrong actions on its own, causing damage even without an attacker.

The magnitude of the challenge:

- [Securing AI agents](#) is the #1 concern for 37% of CISOs surveyed
- 80% of organizations [report risky agent behaviors](#), including improper data exposure and unauthorized system access
- [Only 21% of executives](#) have complete visibility across agent behaviors, permissions, tool usage, or data access

“

AI innovation is accelerating, yet without canonical standards like AIUC-1, the gap between capability and control continues to widen. Closing that gap will require auditable AI security as a baseline, not a differentiator.

**Zachary Elewitz**

Fortune 500 AI Leader, Professor, and Board Member

### Traditional controls won't work

AI agents do not behave like employees. They operate at machine speed, spawn non-human identities, and request access to tools and data across systems that were never designed to coordinate policy in real time. Applying traditional human access models to agents creates standing privilege at scale, and standing privilege becomes standing risk.

“

AI components change constantly across the supply chain, but security controls assume static assets, creating blind spots, friction, and no clear accountability when behavior shifts.

**Omar Khawaja**

VP & Field CISO - Databricks



In production environments, the practical failure mode is not model hallucination. It is over privileged agents holding long lived credentials, chaining across systems, and executing actions without clear ownership or containment boundaries. When a prompt injection or logic flaw occurs, the blast radius is defined by the access the agent already has.

### Identity and access changes:

- Agent usage exacerbates already ineffective identity management approaches - the exponential growth agents introduce requires refocusing on security fundamentals
- Traditional access control policies don't translate effectively to AI agents; organizations are seeing widespread unintended data sharing between business units as agents navigate across systems during task execution
- Agents [spawn non-human identities](#) (NHIs) in security blindspots - API keys, service accounts, authentication tokens - that receive broad, persistent access without human-equivalent safeguards

“


AI agents are not users. They are software that can reason, call tools, and take action at machine speed. If we give agents the same long-lived credentials and broad permissions we designed for humans, we turn convenience into standing access and standing access into standing risk.

**Nancy Wang**

CTO - 1Password, Venture Partner - Felicis Ventures

### The agent challenge as seen in recent incidents:

- [ServiceNow BodySnatcher \(CVE-2025-12420\)](#): Unauthenticated attackers could impersonate administrators and execute AI agents to create backdoor accounts - using only an email address, bypassing MFA and SSO. The attack was particularly notable because A2A implementation allowed admin impersonation.
- [Amazon Q's VS Code extension vulnerability \(CVE-2025-8217\)](#): The inappropriate scope of a GitHub token in the Q Developer CodeBuild configuration enabled a threat actor to commit malicious code directly into the extension's open-source repository, which was then automatically included in a release. A syntax error was the sole factor preventing the code from running.
- [OpenAI plugin ecosystem](#): Supply chain attack compromised agent credentials from 47 enterprise deployments; attackers accessed customer data for months before discovery



### Win conditions for the agent challenge

The companies who will emerge as winners in 2026 recognize that high-level AI governance frameworks fall short in meeting this challenge. Frameworks like NIST AI RMF and ISO 42001 provide valuable organizational structure for AI governance - establishing management systems, risk committees, and documentation requirements. But they weren't designed to answer the technical questions CISOs are asking: How do we prevent prompt injection in production? What access controls should govern agent tool usage? How do we test for cascading failures in multi-agent systems?

The gap is analogous to the difference between ISO 27001 (which defines an information security management system) and SOC 2 (which specifies technical controls). Both are necessary, but only one tells you whether TLS is properly configured, database encryption is enabled, or access logs are being monitored. For AI agents, organizations need frameworks that translate governance principles into specific, auditable technical requirements.

Frameworks like AIUC-1 that are technically grounded fill this gap. Rather than asking "do you have an AI risk committee?", they ask "do your agents validate tool call parameters before execution?" and "can you demonstrate that prompt injection attempts are logged and alerted?" These frameworks bridge the divide between CISO priorities and engineering implementation, providing security teams with specific evidence requirements they can use during vendor evaluations and internal audits.

The result is faster procurement cycles, clearer accountability, and confidence that deployed agents meet enterprise security standards - not just governance checkboxes.

“

AI agents compress the OODA loop (Observe, Orient, Decide, Act) beyond human reaction time, rendering traditional oversight insufficient. Security must be engineered into the agent's decision cycle, not layered on afterward.

**Jonathan Fuller**

CISO - United States Military Academy

## 2. The visibility challenge

**The second challenge facing organizations in 2026 has to do with AI visibility.** It encompasses two interconnected problems:

1. Organizations don't know which AI systems their employees are using - also known as “shadow AI”
2. Even for sanctioned AI deployments, they lack insight into how those systems behave, what data they access, and when their behavior changes.

While the agent challenge focuses on securing AI systems once deployed, the visibility challenge addresses the foundational question: do you even know what AI is operating in your environment?

This isn't just an inventory problem - it's about establishing governance structures that can track AI usage, ownership, and risk across an organization where AI adoption happens through browser tabs and API calls, often with no procurement paper trail.

“

Organizations rarely lose control of AI overnight. Instead, they lose it gradually – when data access isn't tracked and agentic deployments lack proper ownership, allowing the unknown to go undetected.

**Lena Smart**

Ambassador - AIUC-1

### The magnitude of the challenge:

- 1 in 5 organizations [reported breaches due to shadow AI](#) - indeed, 63% of employees [paste company data into personal chatbot accounts](#), including source code, customer data, and internal documents
- Shadow AI breaches [cost \\$670,000 more than standard incidents](#) on average driven by delayed discovery (tools aren't monitored), scope uncertainty (hard to determine what was exposed), and lack of basic controls like DLP or access restrictions
- The average enterprise [has 1,200 unofficial AI applications](#) - 86% are completely blind to their AI data flows

“

Business leaders are pushing AI-adoption, but often without access to the right tools and training, are increasing the risk of shadow AI use.

**Louise McElvogue**

Board Director and Advisor

### Traditional controls won't work

Shadow AI isn't just employees experimenting with ChatGPT. It's a material breach vector operating at scale, outside governance, with access to your most sensitive data. When CISA's acting director uploaded sensitive government contracting documents into public ChatGPT despite multiple security warnings, it wasn't a sophisticated attack - it was the predictable outcome of invisible AI adoption outpacing security controls.

The security playbooks designed to govern software deployment and data access can't keep pace with AI adoption that happens through browser tabs, API calls, and embedding snippets - often without IT approval, procurement review, or any paper trail at all. The organizations trying to get ahead on the challenge, often end up rushing the approval process for AI tools, creating governance debt that needs to be dealt with later.

“

To survive the Agent Challenge of 2026, we must recognise that security is not a reductionist checklist of controls nor a component property to be tested into an agent, but a dynamic control problem to be wired into the organisation at board, management and operational levels.

**Simon Goldsmith**

CISO

**Insights from AIUC-1 consortium members highlight the nuances of the visibility challenge:**

- **The governance gap:** Only 25% of organizations [have comprehensive AI security governance](#) in place. Most companies have developed AI principles or ethics statements, but less than 25% have board-approved AI policies. Business leaders are pushing AI adoption, but often without access to the right tools and training, increasing the risk of shadow AI use. The result is a dangerous bifurcation: in some organizations, adoption is directly stifled due to governance bottlenecks - but in others, it's racing ahead with no governance at all.
- **The "unowned decision" problem:** Small, "non-security" AI decisions in the supply chain - data refreshes, model swaps, prompt changes - quietly introduce real risk with no visibility, audit trail, or ownership. Change control assumes static assets; AI components change constantly, creating blind spots and unclear accountability when behavior shifts. AI security controls deployed as "add-ons" create friction because they're trying to apply traditional controls to a fundamentally new way of operation and emerging non-deterministic threats.

“

AI risk increasingly manifests as unowned operational decisions, not deliberate security failures.

**Gagandeep Singh**

VP, Global Compliance & Certifications - Salesforce

- **The human factor persists:** People remain the weakest link - falling for AI-generated mis- and disinformation, lacking tools or inclination to evaluate whether output is real, useful, or hallucinated. Generative AI data violations [more than doubled last year](#) and shadow AI is preventing business leaders from keeping a lid on sensitive data. The behavioral monitoring required for AI security often feels like surveillance to employees - DLP is a key control, yet still seen as "big brother."

Organizations are struggling to balance security visibility with privacy expectations and trust culture. The risk isn't addressed by the familiar, top-of-mind threats - it requires deliberately ensuring you're addressing the right risks, not falling for availability heuristic.



Having well defined internal and external AI risks and being comfortable operating in gray areas and translating those into actions is crucial to navigate and solve for the challenges ahead.

**Valentina Poghosyan**

Chief Compliance Officer - MongoDB

### The visibility challenge as seen in recent incidents:

- McDonald's AI hiring system leaked millions of candidates' PII through inadequate access controls - security teams didn't know the system existed until the breach
- Major financial services firm discovered 847 unauthorized AI tools after implementing proper discovery - including customer service agents using ChatGPT with full CRM access
- Enterprise SaaS vendor's "AI assistant" feature quietly enabled by default exposed confidential contract terms to LLM training data

### Win conditions for the visibility challenge

The companies who will emerge as winners in 2026 reframe AI security from compliance burden to competitive advantage. Just like seatbelts enabled us to drive faster, proactive AI governance enables organizations to adopt AI with speed and confidence.

This requires a fundamental shift from reactive controls to proactive architecture. Security teams need frameworks that help them say "yes, and here's how" rather than just "no." Industry-led standards provide the structure to move from shadow AI whack-a-mole to systematic governance - giving security and GRC teams the tools to guide their organizations through adoption successfully.



AIUC-1 functions as a centralized, executable control reference layer, normalizing AI governance requirements across threat frameworks, regulations, and best practices to support consistent and automatable implementation.

**Jyoti Wadhwa**

Fmr. Global Head of Product and Cloud Security - NetApp

The winning approach has three components:

- **Visibility first:** Organizations need discovery tools that can identify AI usage across sanctioned and unsanctioned channels - not just endpoint monitoring, but API call patterns, browser extensions, and embedded widgets. This isn't surveillance - it's the foundation for informed decision-making.
- **Structured enablement:** Rather than blanket approvals or bans, organizations can create tiers of AI usage with clear requirements at each level. Low-risk use cases (like code completion without data access) get streamlined approval. High-risk applications (like customer-facing agents with database access) require technical review against specific criteria - ideally mapped to frameworks like AIUC-1 that provide concrete evidence requirements.
- **Continuous governance:** AI systems change constantly. Effective governance means ongoing monitoring of model updates, prompt changes, and tool integrations - not point-in-time assessments. Organizations that build continuous validation into their AI operations avoid both the friction of constant re-review and the risk of drift from approved configurations.

The goal isn't to slow adoption - it's to make secure adoption the path of least resistance.

“

To manage shadow AI while preserving innovation velocity, enterprises must adopt a disciplined operating model that ensures continuous visibility into AI systems and agents, enforces enterprise wide policy, and embeds proactive technical and administrative governance with continuous monitoring across the AI solutions & the agentic ecosystem.

**Adnan Dakhwe**

CISO & CIO - DelphinusCyber

### 3. The trust challenge

**The third challenge we believe will define 2026 is the trust challenge.** You can't trust your inputs. You can't trust your components. The foundations of secure system design - validate inputs, verify dependencies - break down when inputs are natural language and dependencies are opaque neural networks.

Prompt injection moved from research curiosity to production incidents in 2025. OWASP's 2025 LLM Top 10 places prompt injection at #1; 53% of companies [now use RAG and agentic pipelines](#), creating new injection surfaces. The vulnerability exists because LLMs cannot reliably separate instructions from data - inputs affect models even if imperceptible to humans.

“

Model-level safety is necessary but not sufficient. Fine-tuning attacks, extraction attacks, and indirect prompt injection all bypass many model guardrails.

**Professor Sanmi Koyejo**

Leader - Stanford Trustworthy AI Research Lab; Co-founder - VirtueAI

#### The trust challenge as seen in recent incidents:

- Microsoft Copilot "EchoLeak": Zero-click prompt injection [enabled data exfiltration](#) from OneDrive, SharePoint, and Teams without any user interaction
- Amazon Q's VS Code extension: An inappropriately scoped Github token enabled a thread actor to commit malicious code in the extension's open source repository. A syntax error was the sole factor preventing the code from running.
- ServiceNow Now Assist: "Second-order" prompt injection - a low-privilege agent tricked into [asking a higher-privilege agent to perform unauthorized actions](#); initially characterized as "expected behavior"

“

Trust begins with the relationship we have with technical partners. We will never be able to build systems that have no vulnerabilities since the technology changes so fast. Trust must be cultivated continuously through relationship building activities so when the inevitable issue arises, we already have organizational mechanisms to address them.

**Dr. Keri Pearlson**

Senior Lecturer & Principal Research Scientist - MIT Sloan



## Complex AI supply chains compound the trust challenge

The AI supply chain is opaque in ways traditional software supply chains never were. Over 1 million models [were added to Hugging Face](#) in 2024 alone. Malicious models [increased 6.5x year-over-year](#), with new evasion techniques [bypassing standard detection](#).

Traditional supply chain controls can't answer: How do you validate model weight integrity? How do you detect a backdoored pre-trained model? How do you assess training data poisoning? It's almost impossible to create an SBOM when AI is involved - software supply chain management is difficult enough and adding AI makes it nearly impossible.

“

“Non-security” AI decisions in the supply chain, like data refreshes, model swaps, or prompt changes quietly introduce real risk with no visibility, audit trail, or ownership.

**Cassie Crossley**

Fmr. VP, Supply Chain Security - Schneider Electric

## Mapping AI anatomy is key to systemic trust

Trusting an AI system is a fundamental challenge because most organizations operate without a clear understanding of its “anatomy”, treating these systems as “opaque neural networks” rather than transparent architectures. This lack of visibility will erode trust in the core principles of secure systems design.

The components of a modern AI system are not static; it is made up of constantly changing elements across a complex supply chain, including model weights, training data provenance, system prompts and autonomous agents. Additionally, the EU AI Act fundamentally shifts AI transparency from a “best practice” to a strict legal mandate by requiring detailed mapping of an AI system’s internal structure. The act itself introduced dynamic requirements for traceability and observability. This poses both a risk-management and a regulatory-driven challenge.

“

AI risk increasingly manifests as operational decisions made outside formal risk ownership structures. When decision rights, data ownership, and agent permissions are diffuse, small configuration gaps scale into systemic events.

**Julie C. Chatman**

CISO - Human Health Project; CEO - ResilientTech Advisors

## Win conditions for the trust challenge

The companies who will emerge as winners in 2026 require frequent technical testing of deployed agents. Technical safeguards and good governance are sound foundations, but every foundation model update, attack pattern, and system prompt tweak results in different agent behavior.

This demands a shift from "test once, deploy forever" to frequent validation:

- **Assume breach, design for containment:** Since inputs can't be fully trusted, architect systems with the assumption of prompt injection will succeed. Implement least-privilege tool access, segregate high-risk operations, and monitor for anomalous agent behavior patterns.

“

With regard to audit accountability, accountability is deeply important, but so is attribution. While tokens generally prove access, intent can't be inferred. We'll need cryptographic linkage between the human prompt, the agent decision, and the downstream actions - with immutable logs. Accountability should trace to the sponsoring human and system owner. If we can't reconstruct intent, we don't have the proper governance structure in place.

**Peter Holcomb**

Founder & CEO - Optimo IT

- **Red team continuously, not periodically:** The only way to confidently deploy agents when customer trust and brand is on the line is by integrating agent red-teaming as a core tool. Organizations need adversarial testing that evolves with new attack patterns - not annual penetration tests.
- **Demand supply chain transparency:** Vendor due diligence must address upstream and downstream responsibilities - including monitoring for models being updated without customer knowledge, documenting training data provenance, and providing evidence of component validation processes.

“

“Inside out” is the only way to handle AI systems, emphasizing that security must be architected into the system's core rather than applied at the perimeter.

**Rohit Parchuri**

CISO and Advisor

# Priorities for organizations adopting AI in 2026

The companies that will win in 2026 are those that recognize that the era of vibe adoption is over. They will reframe AI governance from a compliance burden to a competitive advantage.

The path forward requires addressing three interconnected challenges with technically grounded approaches:

## Win conditions for the agent challenge

- Adopt technically grounded frameworks that answer specific security questions (prompt injection prevention, tool call validation, cascading failure detection) rather than high-level governance principles
- Bridge the gap between CISO priorities and engineering implementation with auditable technical requirements
- Move beyond "do you have an AI risk committee?" to "can you demonstrate that prompt injection attempts are logged and alerted?"

“

AI security controls as an “add-on” inherently create friction as we are trying to apply traditional controls to a new way of operation. Understanding how the introduction of AI into an environment changes data flows, processes, and the overall risk landscape will help us ensure we understand the true impact of AI, allowing us to more effectively implement controls.

**Mandy Andress**

CISO - Elastic

## Win conditions for the visibility challenge

- Implement discovery tools that identify AI usage across sanctioned and unsanctioned channels - establishing visibility as the foundation for informed decision-making
- Create structured enablement with tiered AI usage policies that make secure adoption the path of least resistance
- Build continuous governance that monitors model updates, prompt changes, and tool integrations - not just point-in-time assessments

## Win Conditions for 3 Challenges

| The <b>Agent</b> Challenge       | The <b>Trust</b> Challenge       | The <b>Visibility</b> Challenge   |
|----------------------------------|----------------------------------|-----------------------------------|
| Technically grounded frameworks  | Implement AI discovery tools     | Assume breach, design containment |
| Bridge CISO-engineering gap      | Tiered usage policies            | Frequent evolving red teaming     |
| Auditable technical requirements | Continuous governance monitoring | Demand supply chain transparency  |

### Win conditions for the trust challenge

- Assume breach and design for containment with least-privilege tool access, operation segregation, and anomalous behavior monitoring
- Integrate frequent red teaming that evolves with new attack patterns
- Demand supply chain transparency from vendors, including model update notifications, training data provenance, and component validation evidence

“

As AI's use for many businesses moves from experimentation and into execution, the risk profile is changing for many. As business and technology leaders, we need standards that translate intent into testable technical controls, and which complement existing frameworks to help fill the gap.

**Craig Weatherhead**

SVP IT Infrastructure & Security - Fastenal

The pressure to adopt AI is real, and the potential benefits are immense. But success in 2026 belongs to organizations that refuse to choose between speed and security. Industry-led frameworks like AIUC-1 provide the technical specificity that security teams need to evaluate vendors systematically, guide internal deployments confidently, and prove to boards, regulators, and customers that their AI systems meet enterprise standards.

**The question is no longer whether to adopt AI - it's whether you have the frameworks, processes, and visibility to adopt it safely.**

# AIUC-1

[aiuc-1.com](http://aiuc-1.com) | [contact@aiuc.com](mailto:contact@aiuc.com)